

Fitting time series models

ARIMA MODELS IN PYTHON



James Fulton

Climate informatics researcher

Creating a model

```
from statsmodels.tsa.arima.model import ARIMA
```

```
# This is an ARMA(p,q) model  
model = ARIMA(timeseries, order=(p, 0, q))
```

Creating AR and MA models

```
ar_model = ARIMA(timeseries, order=(p,0,0))
```

```
ma_model = ARIMA(timeseries, order=(0,0,q))
```

Fitting the model and fit summary

```
model = ARIMA(timeseries, order=(2,0,1))  
results = model.fit()
```

```
print(results.summary())
```

Fit summary

SARIMAX Results

=====

Dep. Variable:

y

No. Observations:

1000

Model:

ARMA(2, 1)

Log Likelihood

148.580

Date:

Thu, 25 Apr 2022

AIC

-287.159

Time:

22:57:00

BIC

-262.621

Sample:

0

HQIC

-277.833

Covariance Type:

opg

=====

coef

std err

z

P>|z|

[0.025

0.975]

const

-0.0017

0.012

-0.147

0.883

-0.025

0.021

ar.L1.y

0.5253

0.054

9.807

0.000

0.420

0.630

ar.L2.y

-0.2909

0.042

-6.850

0.000

-0.374

-0.208

ma.L1.y

0.3679

0.052

7.100

0.000

0.266

0.469

Fit summary

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:          1000
Model:                ARMA(2, 1)  Log Likelihood          148.580
Date:                Thu, 25 Apr 2022    AIC                -287.159
Time:                22:57:00    BIC                -262.621
Sample:                0      HQIC                -277.833
Covariance Type:      opg
```

Fit summary

	coef	std err	z	P> z	[0.025	0.975]
const	-0.0017	0.012	-0.147	0.883	-0.025	0.021
ar.L1.y	0.5253	0.054	9.807	0.000	0.420	0.630
ar.L2.y	-0.2909	0.042	-6.850	0.000	-0.374	-0.208
ma.L1.y	0.3679	0.052	7.100	0.000	0.266	0.469
sigma2	1.6306	0.339	6.938	0.000	0.583	1.943

Introduction to ARMAX models

- Exogenous ARMA
- Use external variables as well as time series
- $\text{ARMAX} = \text{ARMA} + \text{linear regression}$

ARMAX equation

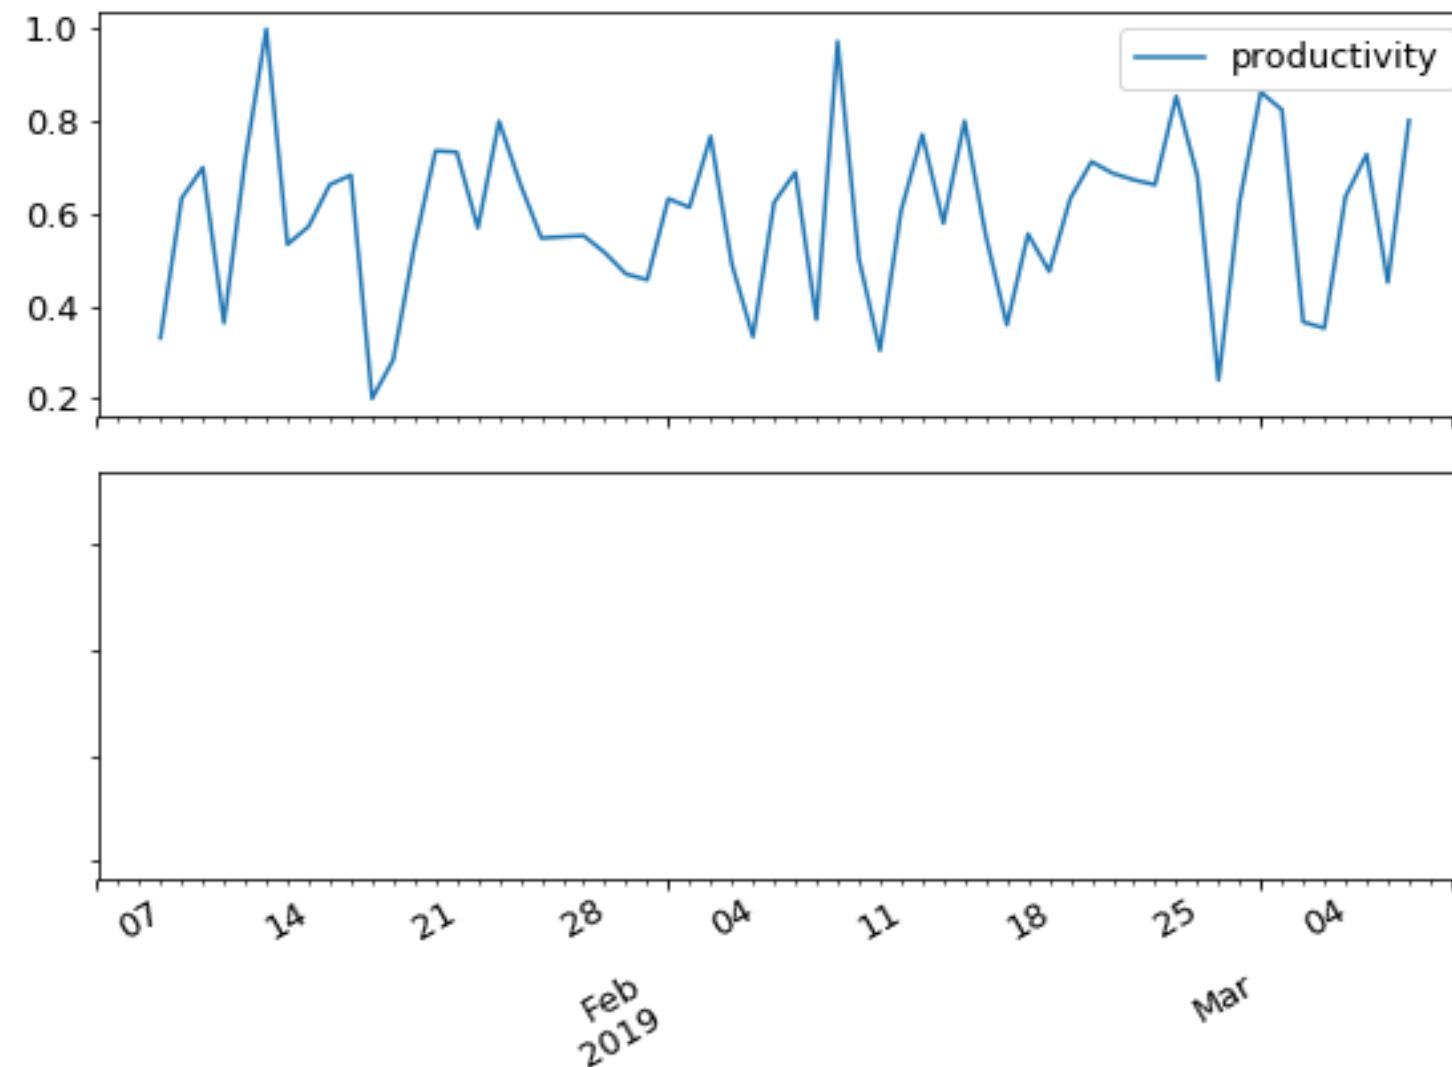
ARMA(1,1) model :

$$y_t = a_1 y_{t-1} + m_1 \epsilon_{t-1} + \epsilon_t$$

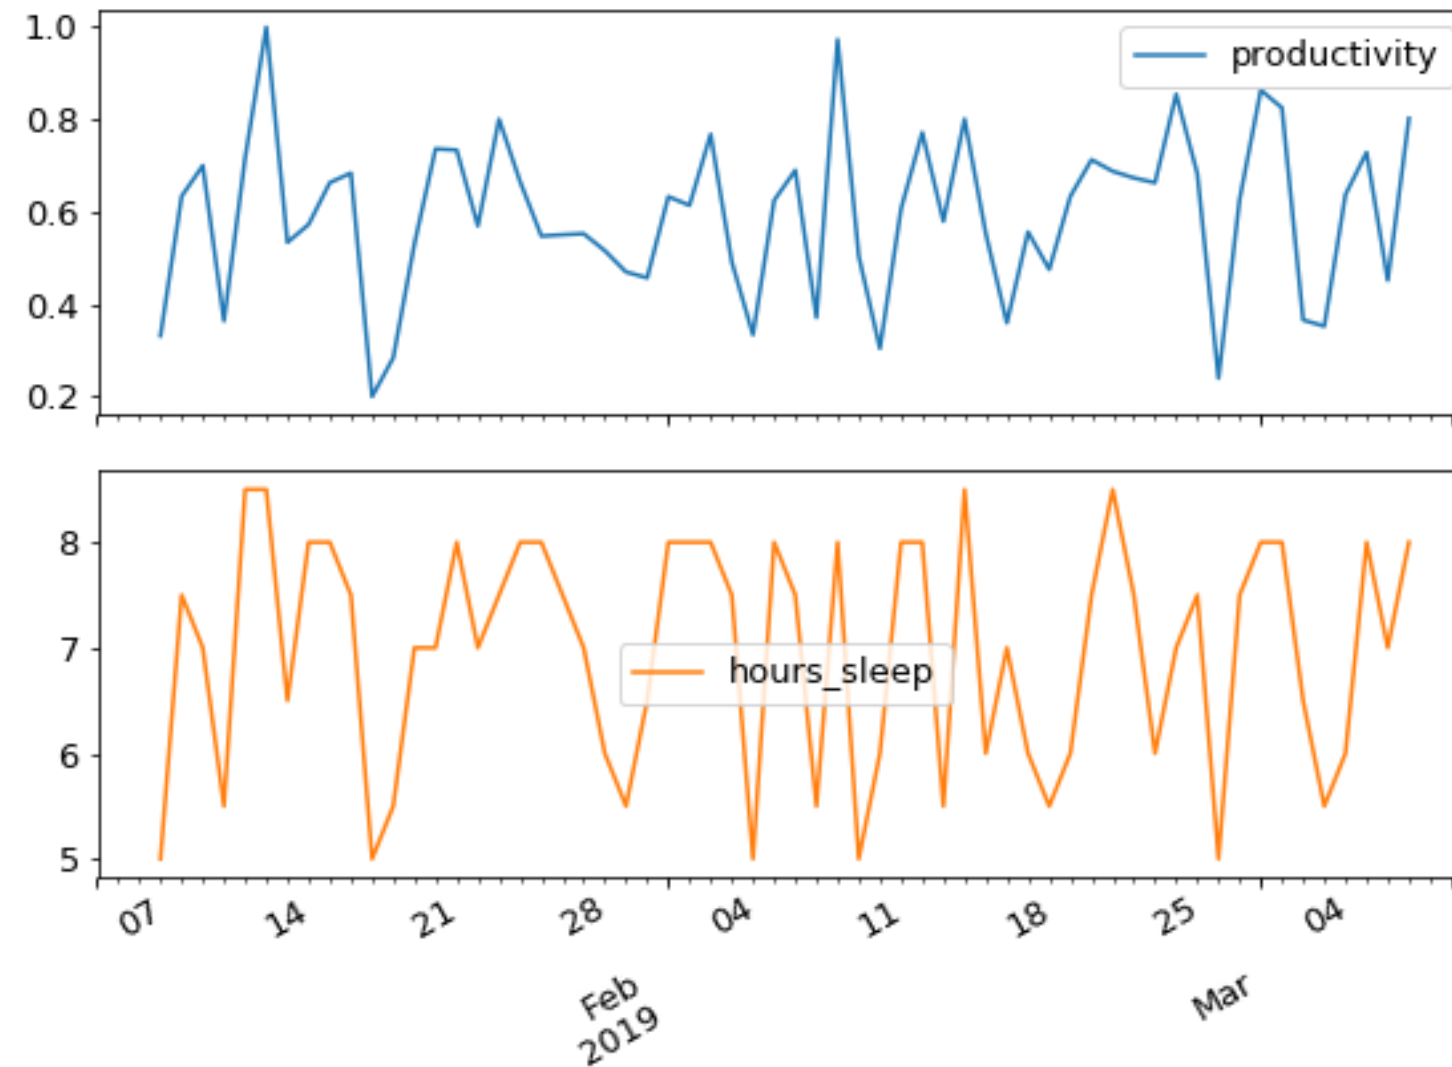
ARMAX(1,1) model :

$$y_t = x_1 z_t + a_1 y_{t-1} + m_1 \epsilon_{t-1} + \epsilon_t$$

ARMAX example



ARMAX example



Fitting ARMAX

```
# Instantiate the model
model = ARIMA(df['productivity'], order=(2,0,1), exog=df['hours_sleep'])

# Fit the model
results = model.fit()
```

ARMAX summary

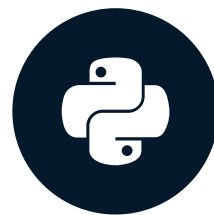
	coef	std err	z	P> z	[0.025	0.975]
const	-0.1936	0.092	-2.098	0.041	-0.375	-0.013
x1	0.1131	0.013	8.602	0.000	0.087	0.139
ar.L1.y	0.1917	0.252	0.760	0.450	-0.302	0.686
ar.L2.y	-0.3740	0.121	-3.079	0.003	-0.612	-0.136
ma.L1.y	-0.0740	0.259	-0.286	0.776	-0.581	0.433

Let's practice!

ARIMA MODELS IN PYTHON

Forecasting

ARIMA MODELS IN PYTHON



James Fulton

Climate informatics researcher

Predicting the next value

Take an AR(1) model

$$y_t = a_1 y_{t-1} + \epsilon_t$$

Predict next value

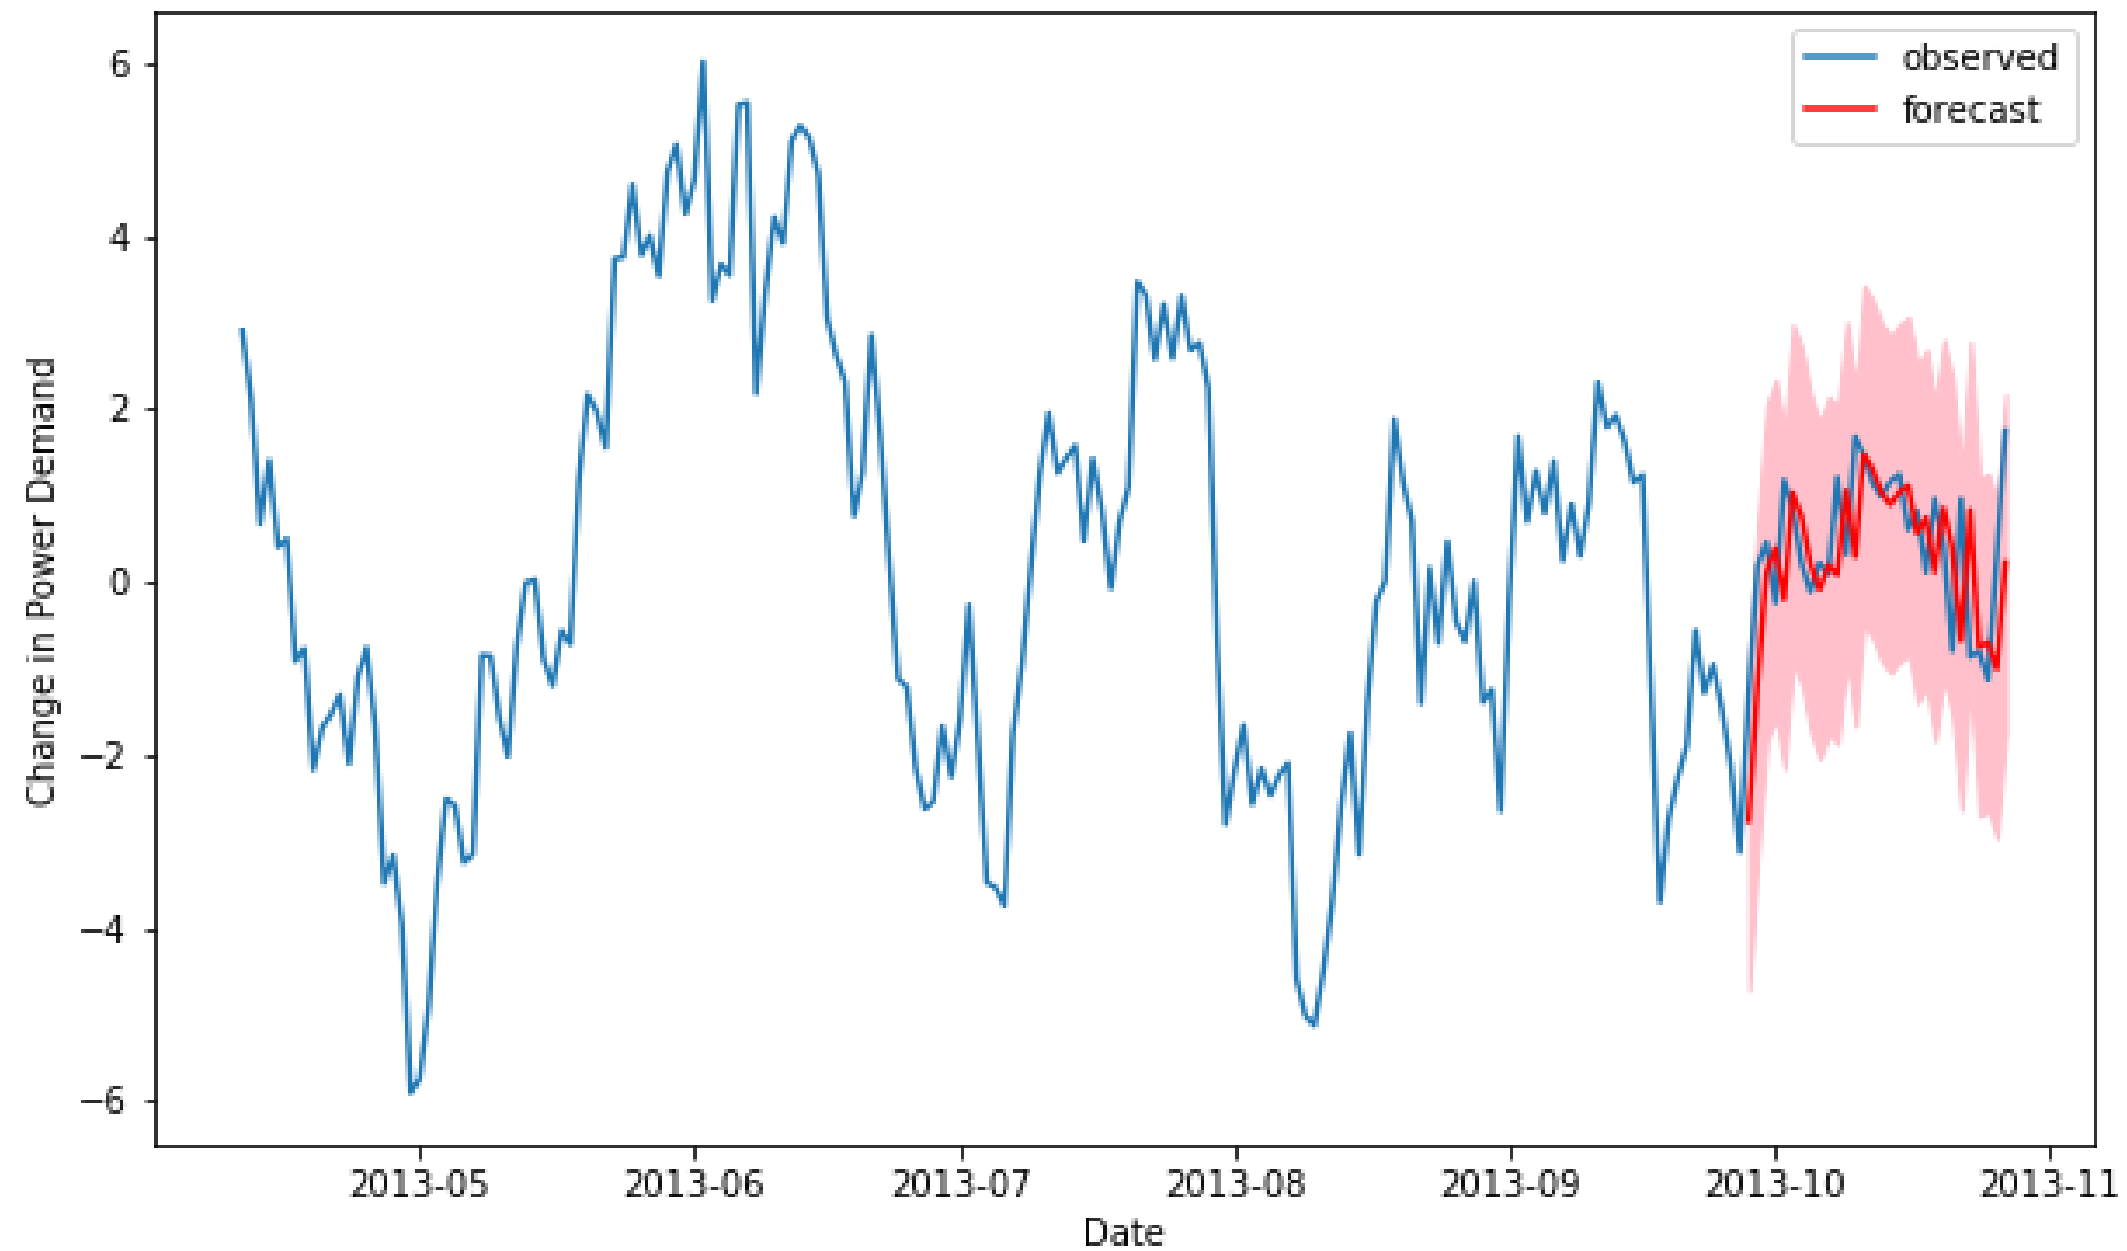
$$y_t = 0.6 \times 10 + \epsilon_t$$

$$y_t = 6.0 + \epsilon_t$$

Uncertainty on prediction

$$5.0 < y_t < 7.0$$

One-step-ahead predictions



Making one-step-ahead predictions

```
# Make predictions for last 25 values
results = model.fit()
# Make in-sample prediction
forecast = results.get_prediction(start=-25)
```

Making one-step-ahead predictions

```
# Make predictions for last 25 values
results = model.fit()
# Make in-sample prediction
forecast = results.get_prediction(start=-25)
# forecast mean
mean_forecast = forecast.predicted_mean
```

Predicted mean is a pandas series

2013-10-28	1.519368
2013-10-29	1.351082
2013-10-30	1.218016

Confidence intervals

```
# Get confidence intervals of forecasts  
confidence_intervals = forecast.conf_int()
```

Confidence interval method returns `pandas` DataFrame

	lower y	upper y
2013-09-28	-4.720471	-0.815384
2013-09-29	-5.069875	0.112505
2013-09-30	-5.232837	0.766300
2013-10-01	-5.305814	1.282935
2013-10-02	-5.326956	1.703974

Plotting predictions

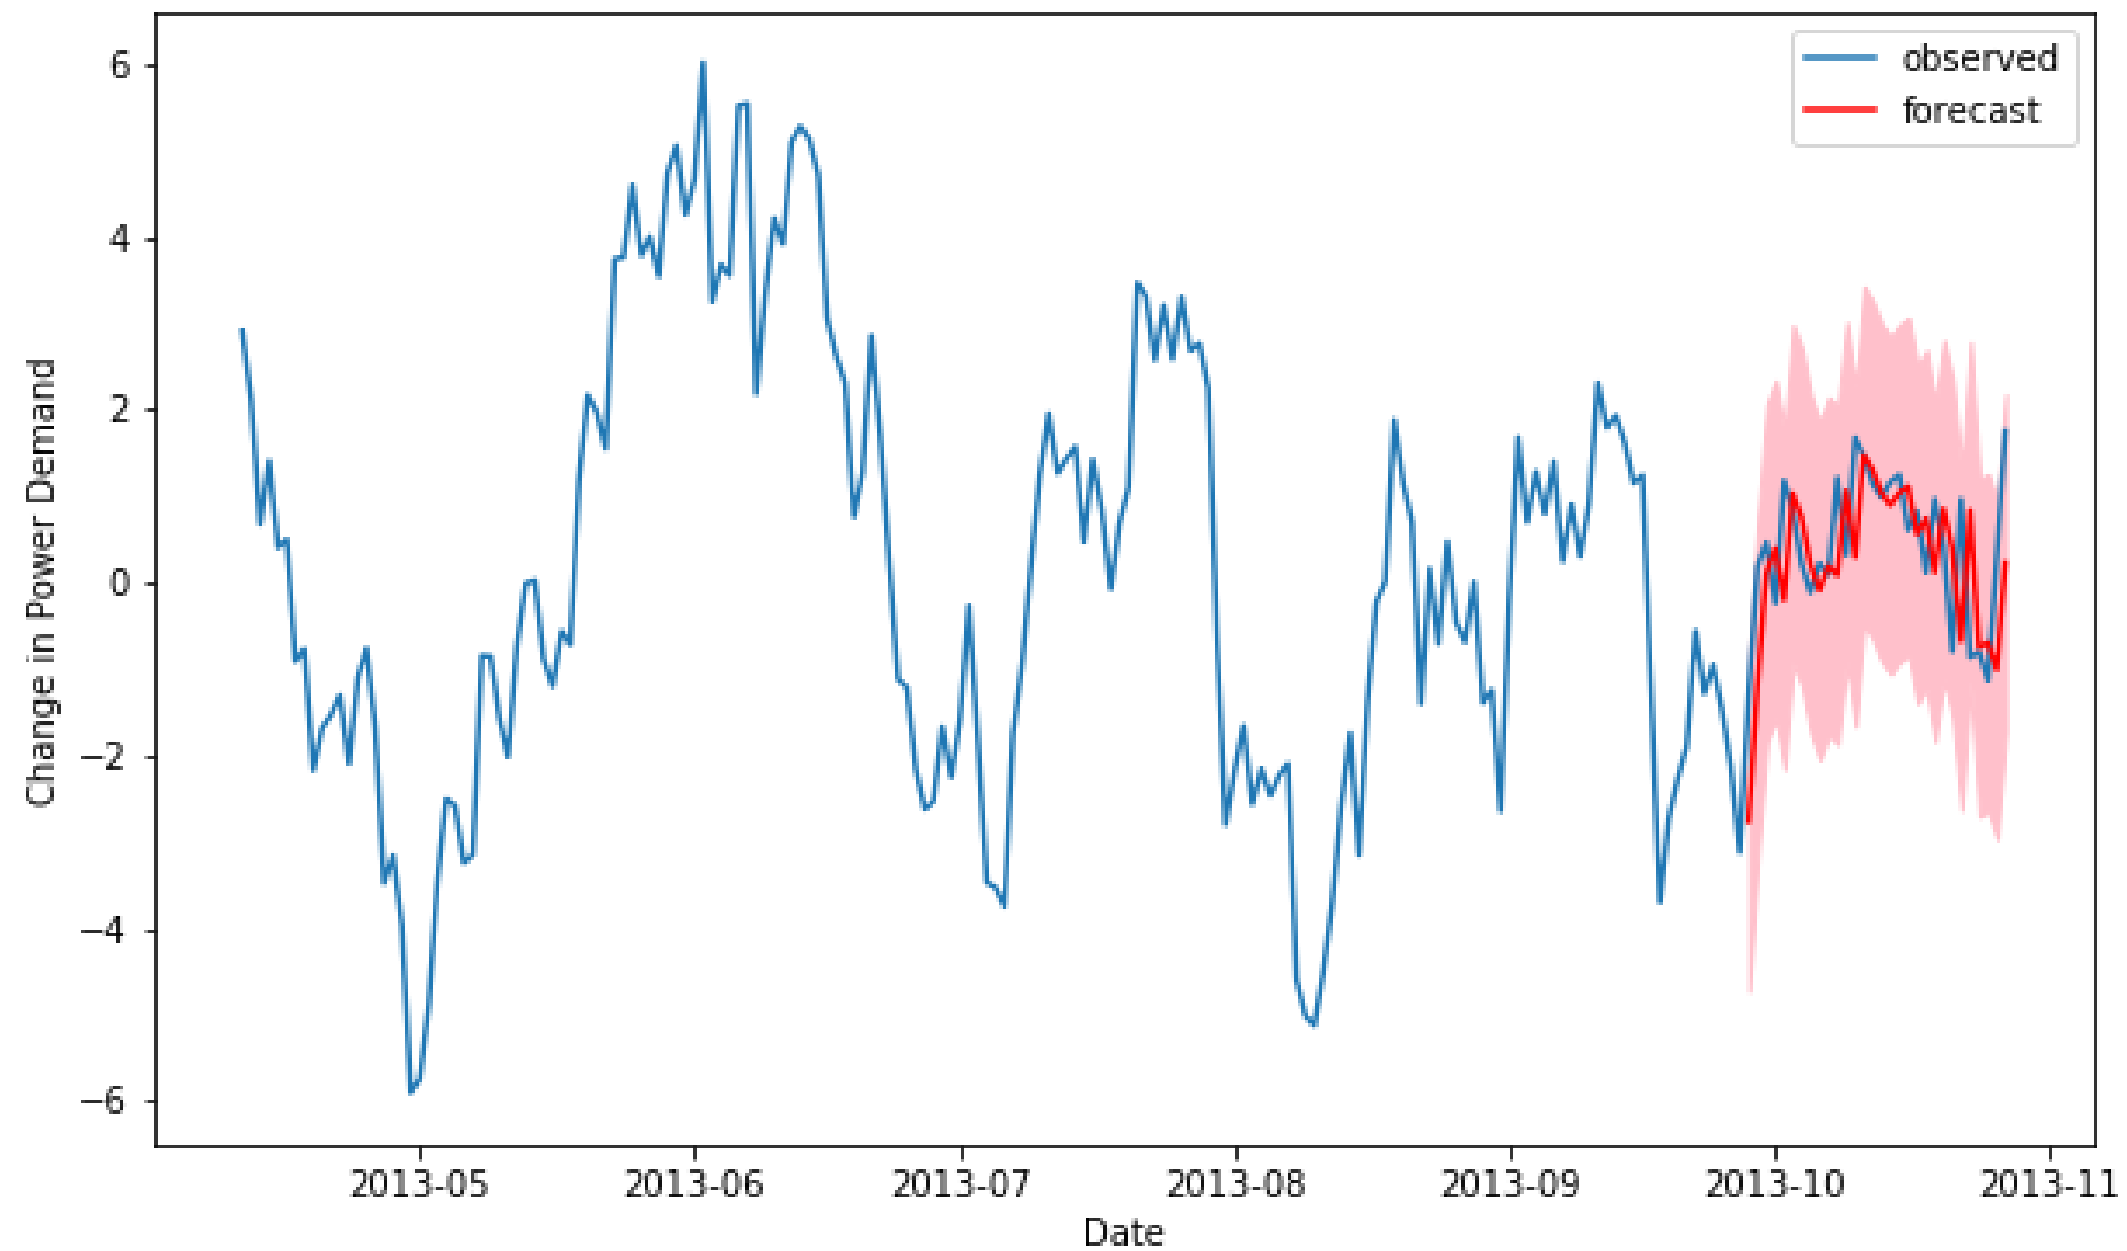
```
plt.figure()

# Plot prediction
plt.plot(dates,
         mean_forecast.values,
         color='red',
         label='forecast')

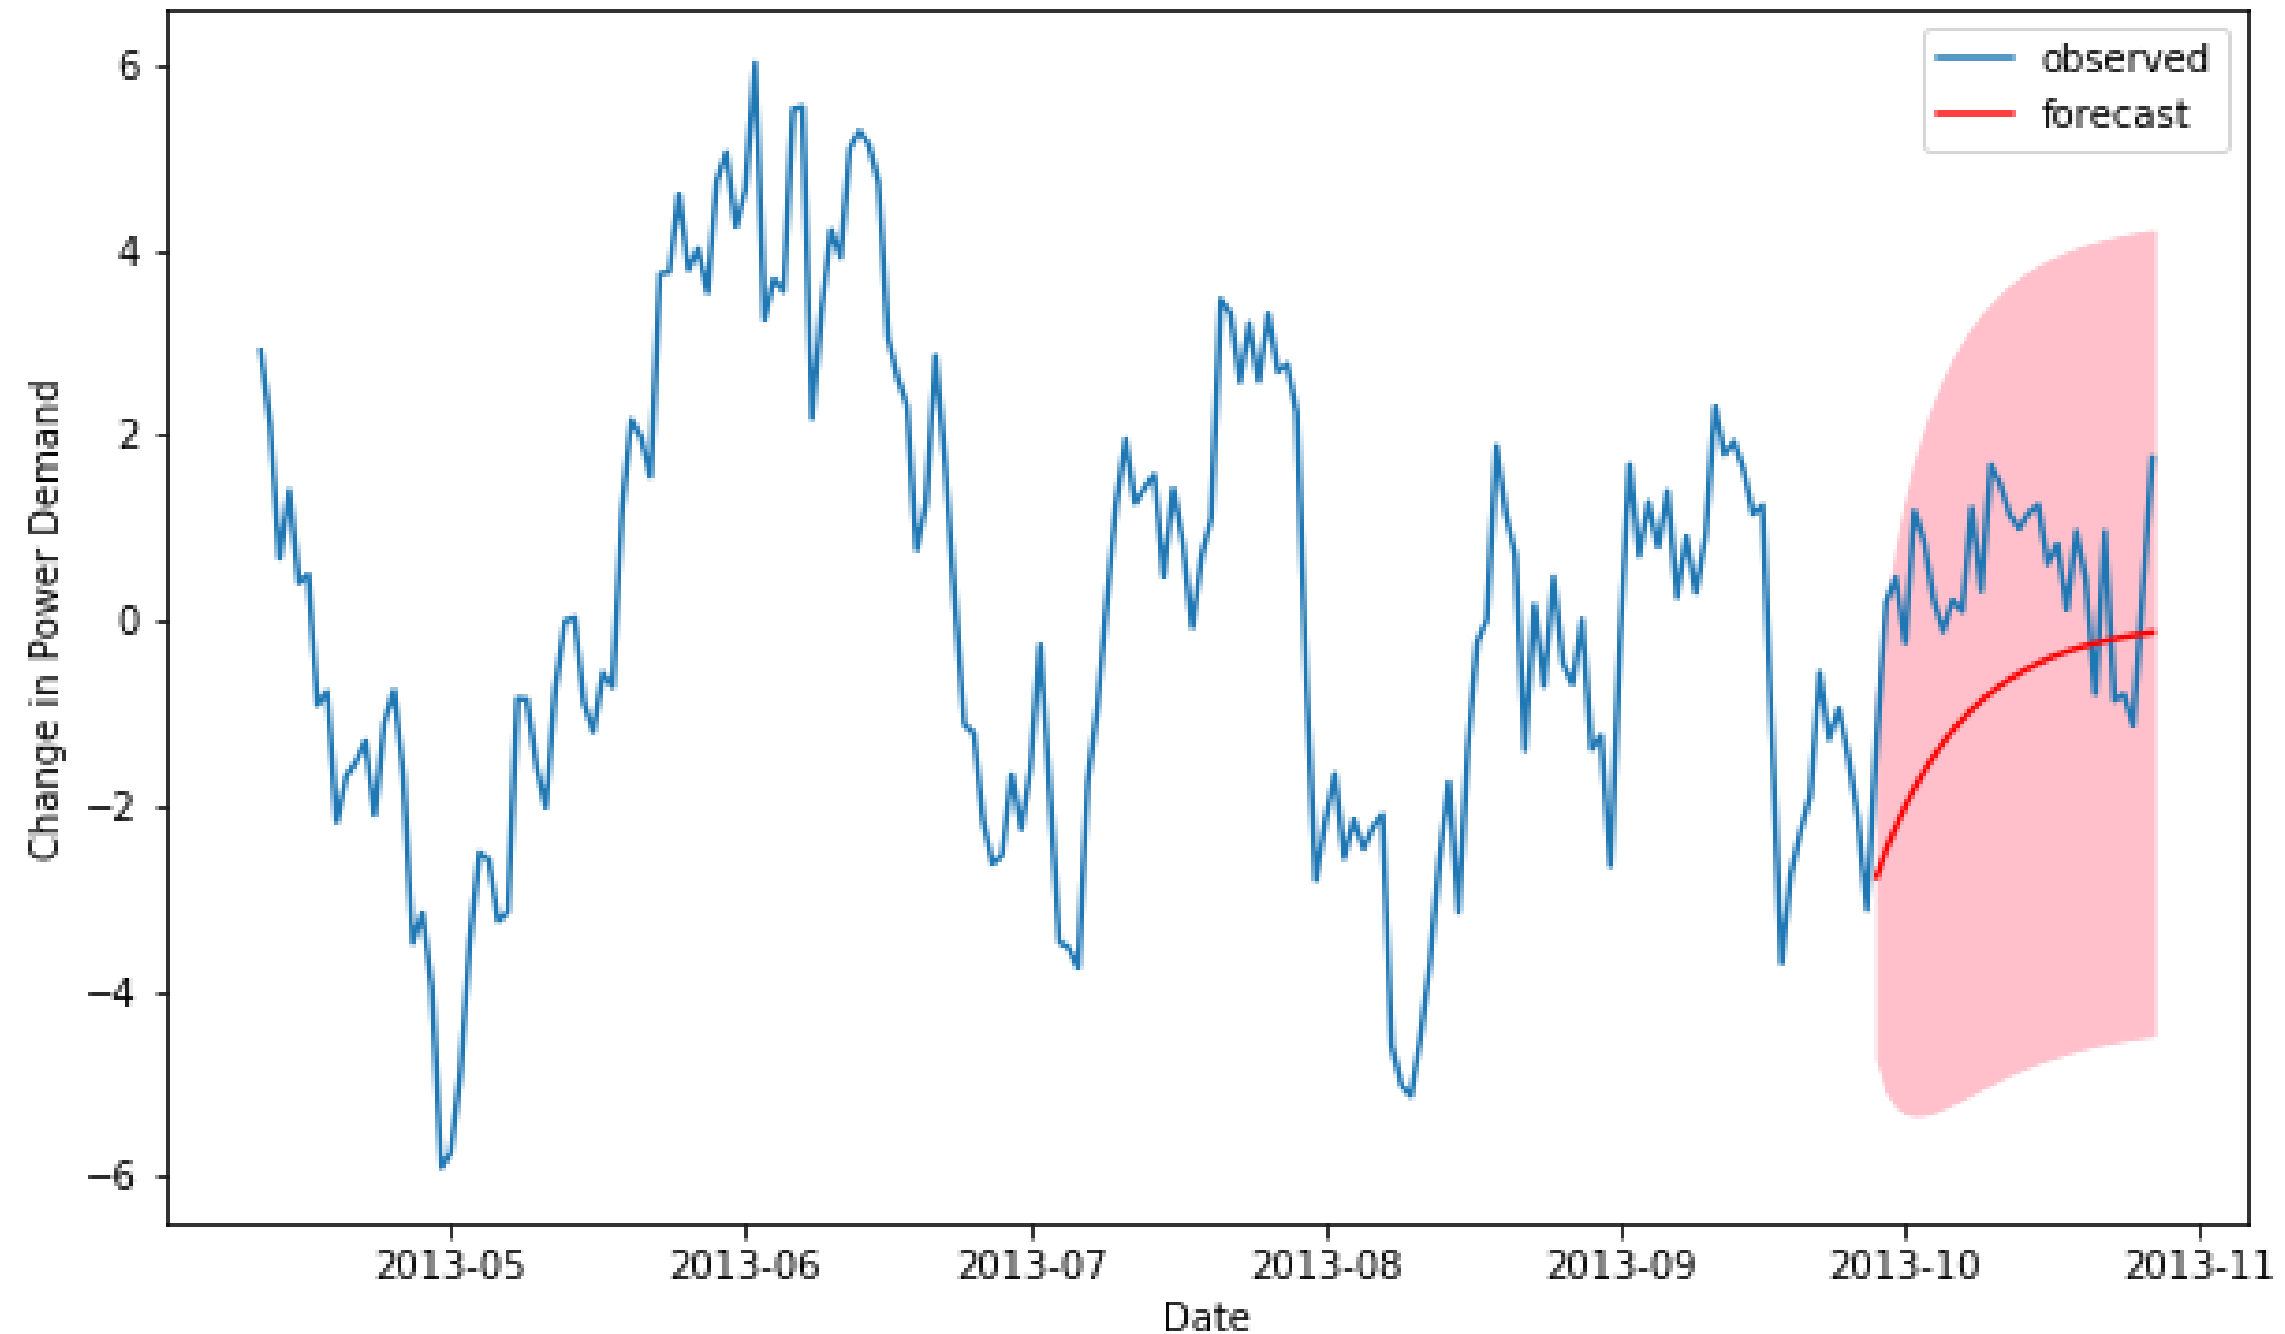
# Shade uncertainty area
plt.fill_between(dates, lower_limits, upper_limits, color='pink')

plt.show()
```

Plotting predictions



Dynamic predictions



Making dynamic predictions

```
results = model.fit()  
forecast = results.get_prediction(start=-25, dynamic=True)
```

```
# forecast mean  
mean_forecast = forecast.predicted_mean  
  
# Get confidence intervals of forecasts  
confidence_intervals = forecast.conf_int()
```


Forecasting out of sample

```
forecast = results.get_forecast(steps=20)
```

```
# forecast mean
```

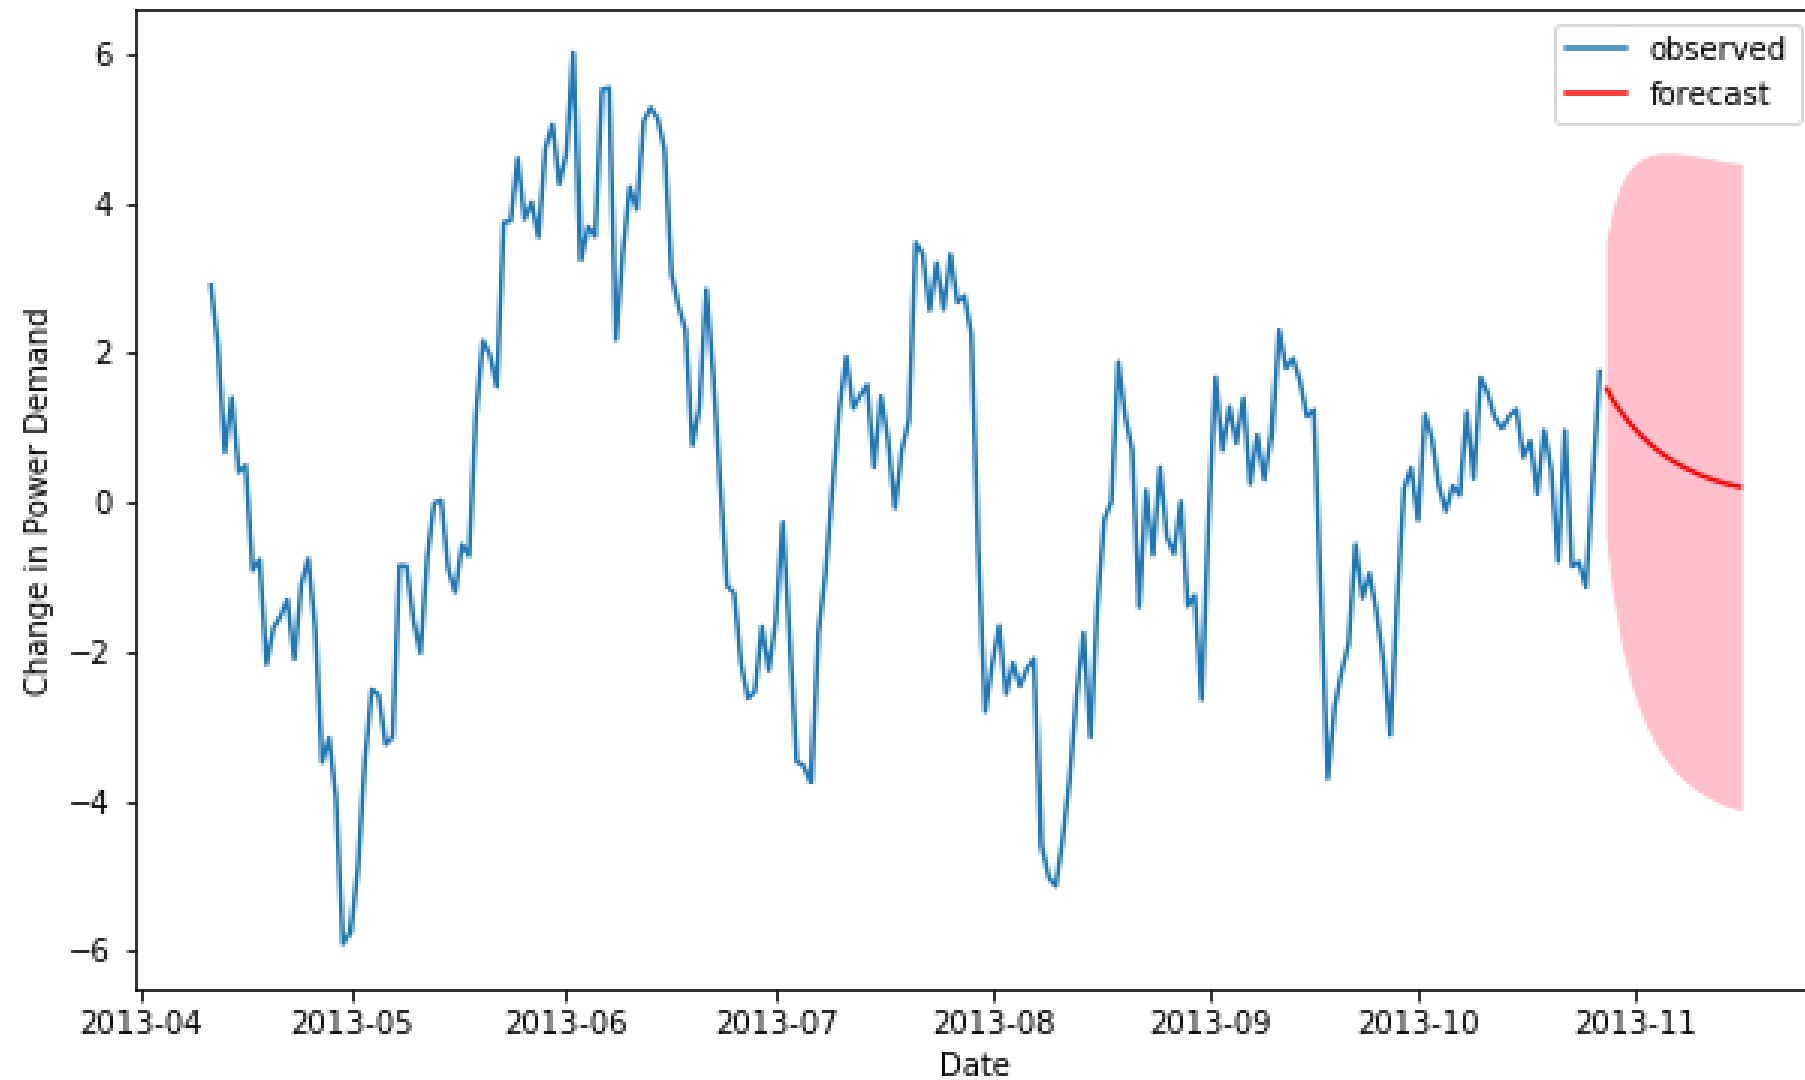
```
mean_forecast = forecast.predicted_mean
```

```
# Get confidence intervals of forecasts
```

```
confidence_intervals = forecast.conf_int()
```

Forecasting out of sample

```
forecast = results.get_forecast(steps=20)
```



Let's practice!
ARIMA MODELS IN PYTHON

Introduction to ARIMA models

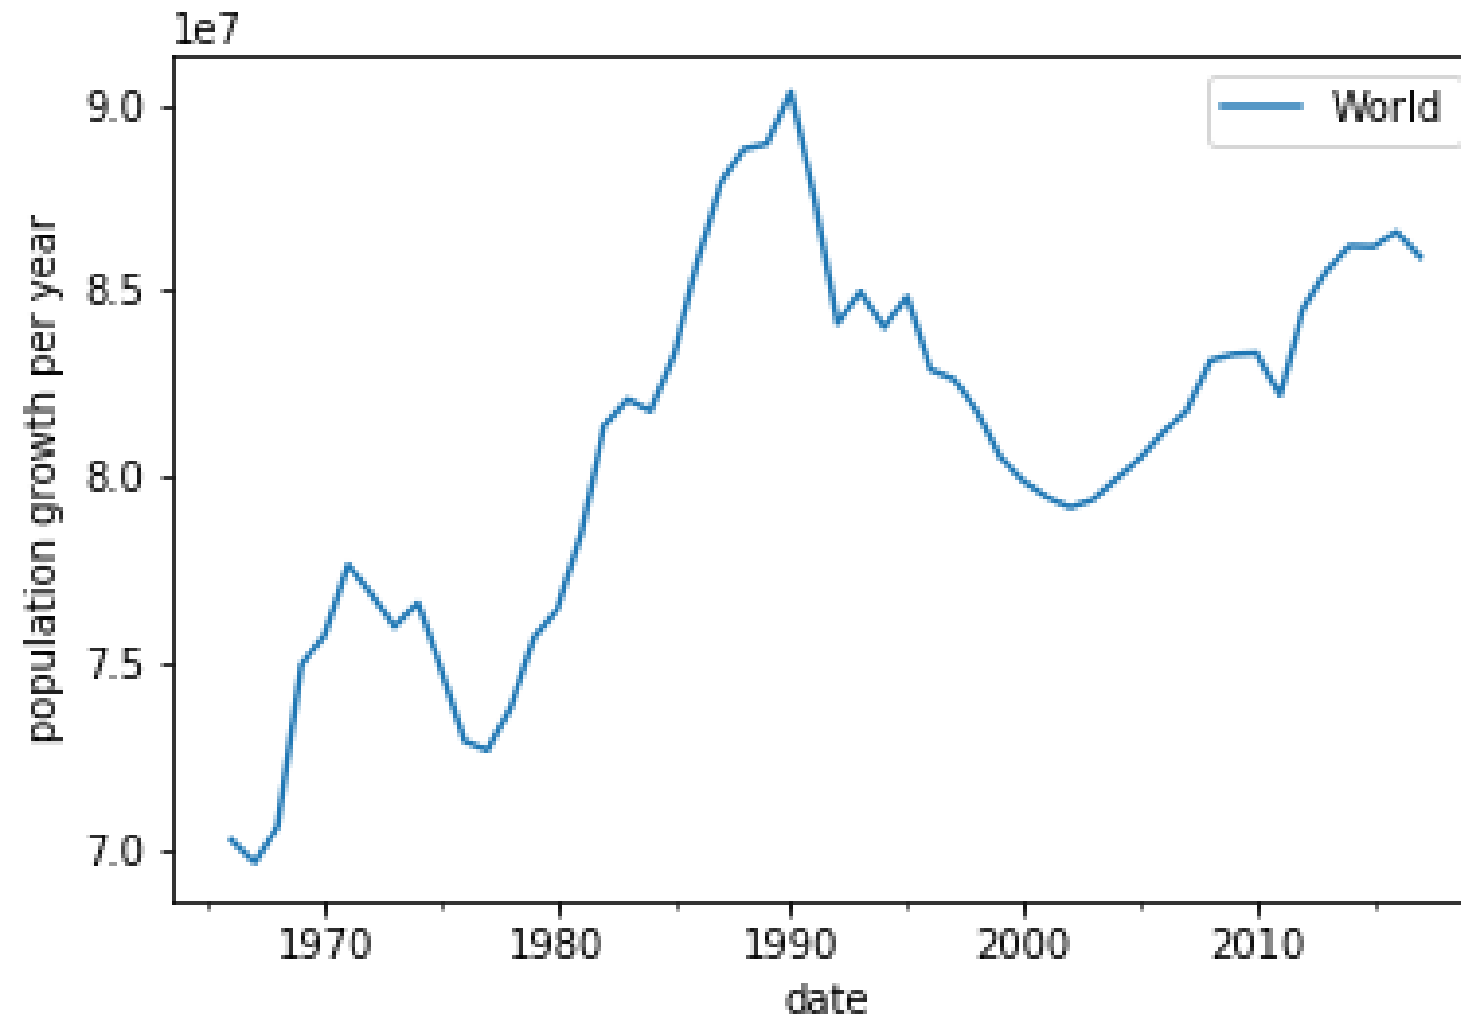
ARIMA MODELS IN PYTHON



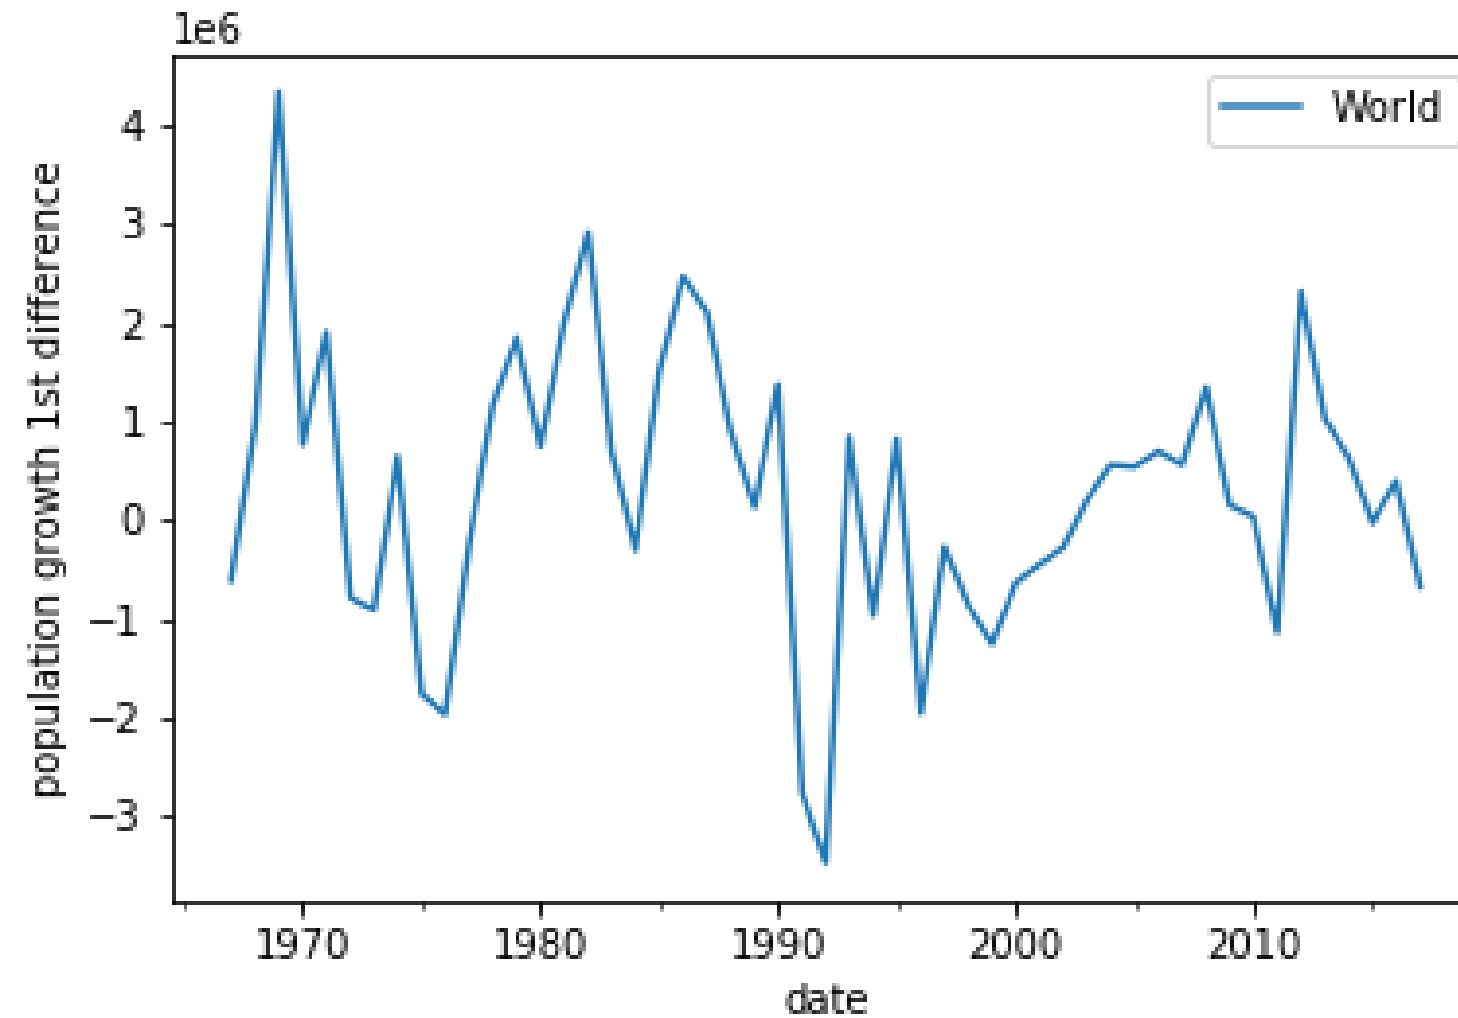
James Fulton

Climate informatics researcher

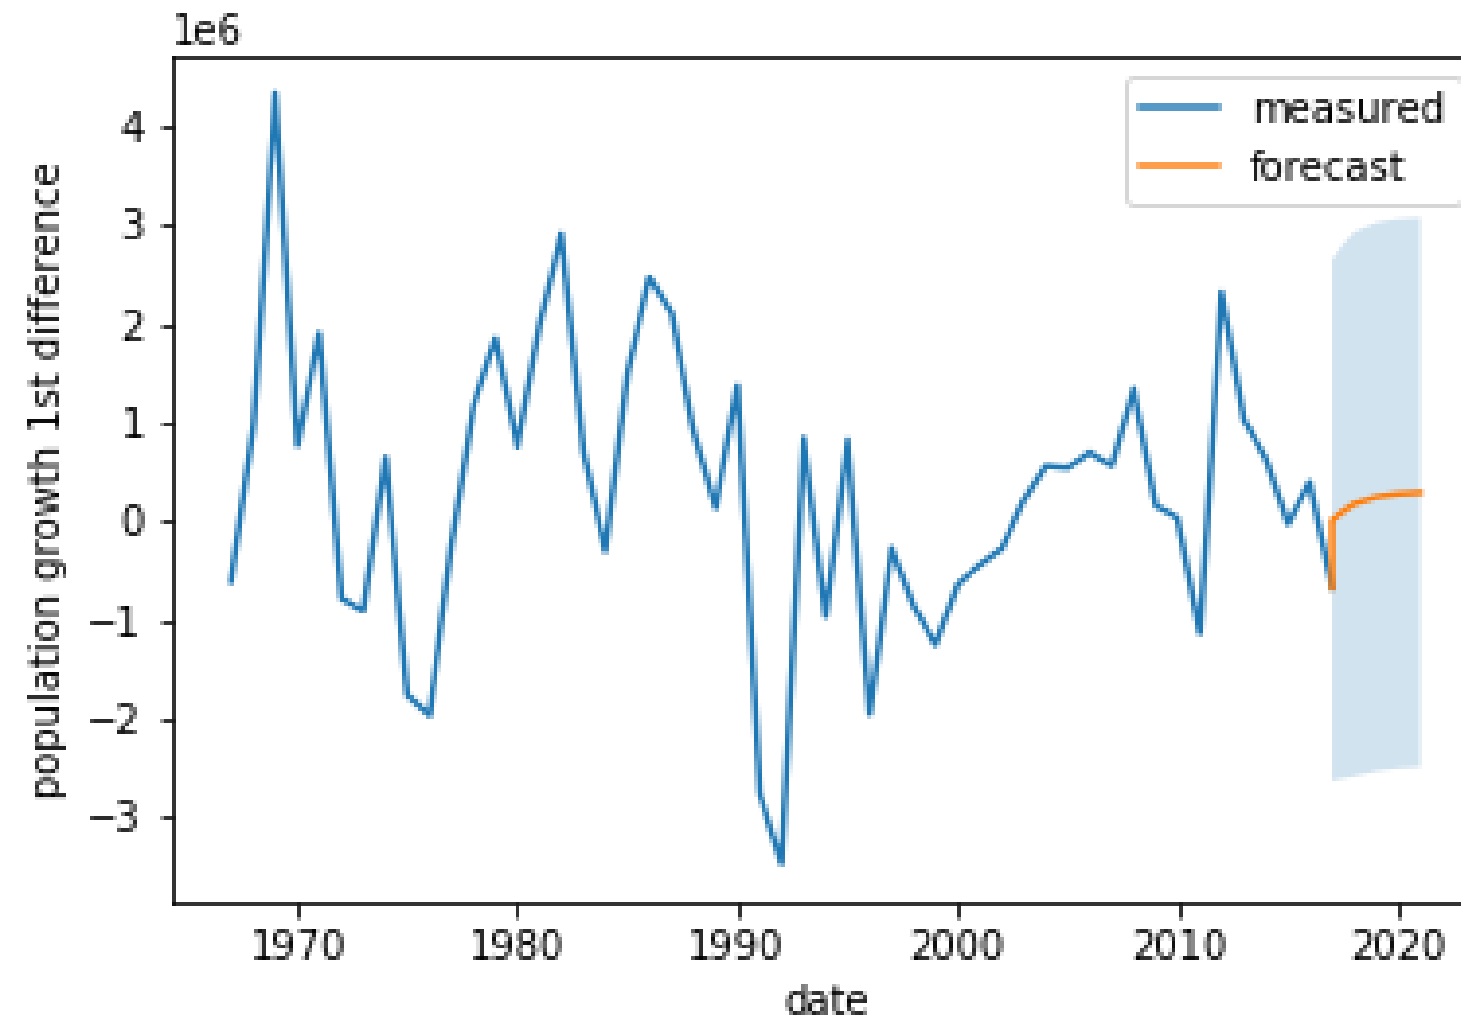
Non-stationary time series recap



Non-stationary time series recap



Forecast of differenced time series



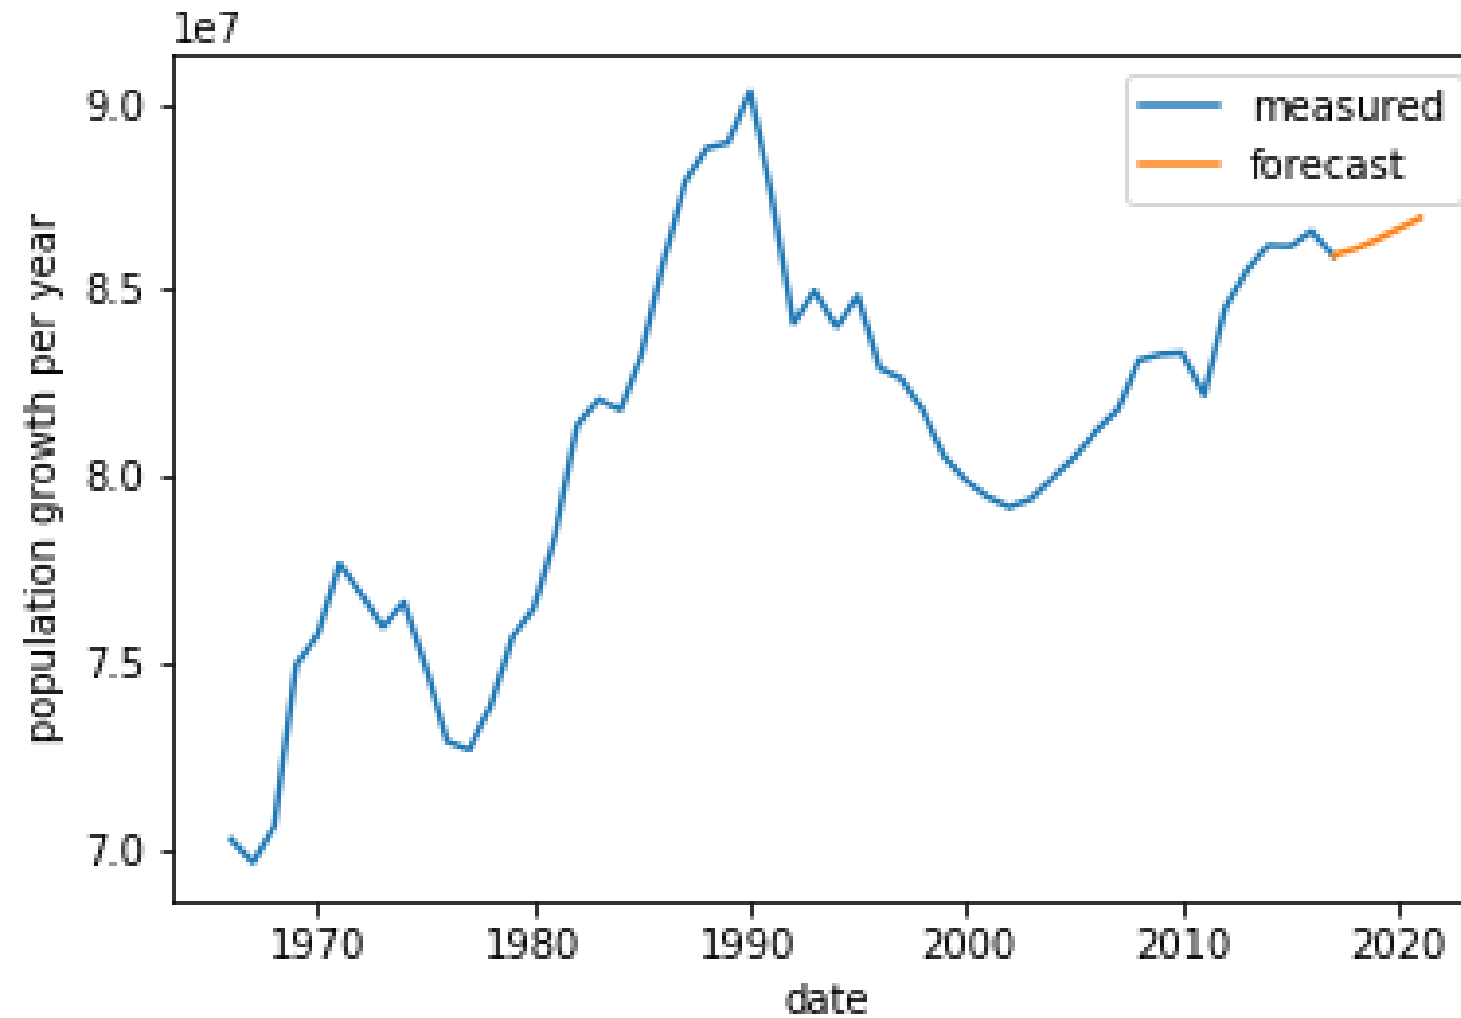
Reconstructing original time series after differencing

```
diff_forecast = results.get_forecast(steps=10).predicted_mean  
from numpy import cumsum  
mean_forecast = cumsum(diff_forecast)
```


Reconstructing original time series after differencing

```
diff_forecast = results.get_forecast(steps=10).predicted_mean  
from numpy import cumsum  
mean_forecast = cumsum(diff_forecast) + df.iloc[-1,0]
```

Reconstructing original time series after differencing



The ARIMA model

- Take the difference
- Fit ARMA model
- Integrate forecast

Can we avoid doing so much work?

Yes!

ARIMA - Autoregressive Integrated Moving Average

Using the ARIMA model

```
from statsmodels.tsa.arima.model import ARIMA  
model = ARIMA(df, order=(p,d,q))
```

- p - number of autoregressive lags
- d - order of differencing
- q - number of moving average lags

$$\text{ARIMA}(p, 0, q) = \text{ARMA}(p, q)$$

Using the ARIMA model

```
# Create model
model = ARIMA(df, order=(2,1,1))

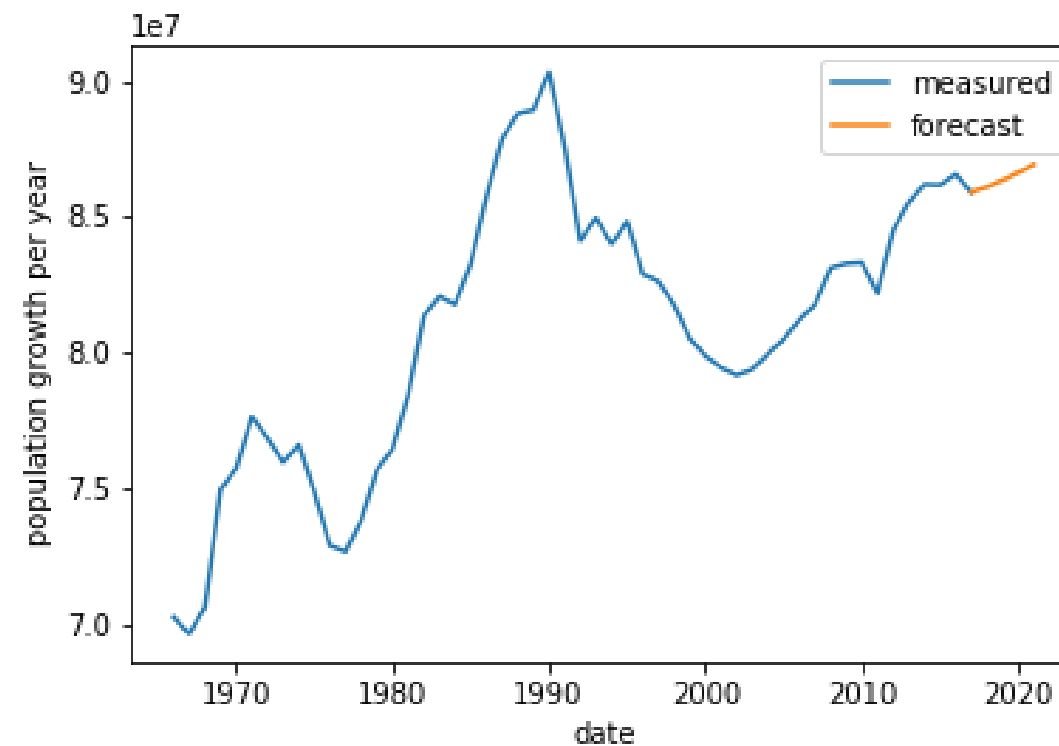
# Fit model
model.fit()

# Make forecast
mean_forecast = results.get_forecast(steps=10).predicted_mean
```

Using the ARIMA model

```
# Make forecast
```

```
mean_forecast = results.get_forecast(steps=steps).predicted_mean
```



Picking the difference order

```
adf = adfuller(df.iloc[:,0])  
print('ADF Statistic:', adf[0])  
print('p-value:', adf[1])
```

```
ADF Statistic: -2.674  
p-value: 0.0784
```

```
adf = adfuller(df.diff().dropna().iloc[:,0])  
print('ADF Statistic:', adf[0])  
print('p-value:', adf[1])
```

```
ADF Statistic: -4.978  
p-value: 2.44e-05
```

Picking the difference order

```
model = ARIMA(df, order=(p, 1, q))
```


Let's practice!
ARIMA MODELS IN PYTHON